

sms-spam-v1-stop

April 23, 2019

V1: + Delete the stop words + All models uses the TfidfVectorizer to do the the preprocessing
Goal of this notebook to test several classifiers on the data set with different features

0.0.1 Let's begin

First of all necessary imports

```
In [1]: import numpy as np
import pandas as pd
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sns
import string
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from nltk.corpus import stopwords
from sklearn.preprocessing import LabelEncoder
%matplotlib inline
```

Let's read the data from csv file

```
In [2]: sms = pd.read_csv('../data/sms-spam.csv',delimiter=',',encoding='latin-1')

sms.head()
```

```
Out[2]:
```

	v1	v2	Unnamed: 2	\
0	ham	Go until jurong point, crazy.. Available only ...		NaN
1	ham	Ok lar... Joking wif u oni...		NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...		NaN
3	ham	U dun say so early hor... U c already then say...		NaN
4	ham	Nah I don't think he goes to usf, he lives aro...		NaN

	Unnamed: 3	Unnamed: 4
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

Now drop “unnamed” columns and rename v1 and v2 to “label” and “message”

```
In [3]: sms = sms.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], axis=1)
        sms = sms.rename(columns = {'v1': 'label', 'v2': 'message'})
```

Let’s look into our data

```
In [4]: sms.groupby('label').describe()
```

```
Out[4]:
```

	message	count	unique	top freq
label				
ham		4825	4516	Sorry, I'll call later 30
spam		747	653	Please call our customer service representativ... 4

Intresting that “Sorry, I’ll call later” appears only 30 times here =)

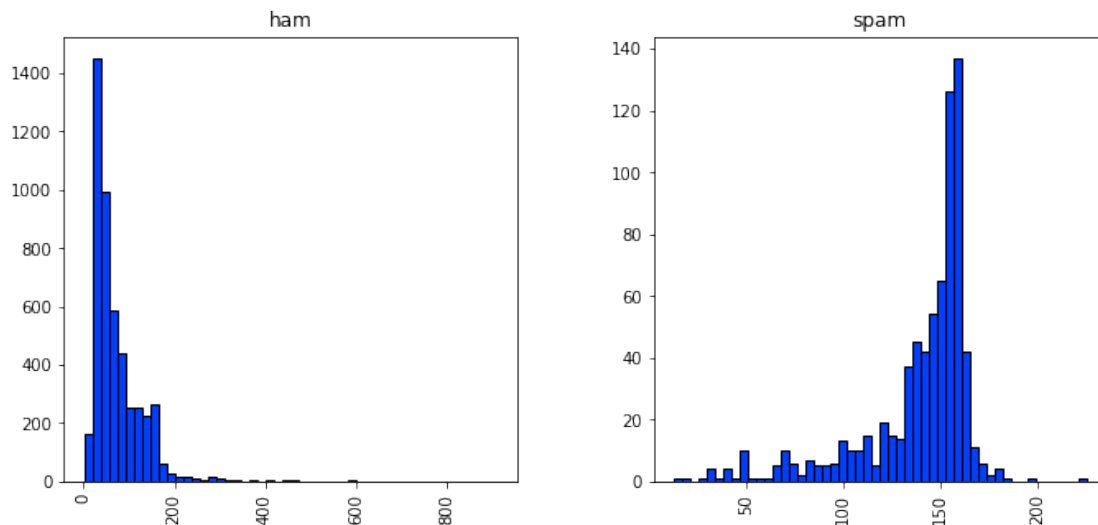
Now let’s create new feature “message length” and plot it to see if it’s of any interest

```
In [5]: sms['length'] = sms['message'].apply(len)
        sms.head()
```

```
Out[5]:
```

	label	message	length
0	ham	Go until jurong point, crazy.. Available only ...	111
1	ham	Ok lar... Joking wif u oni...	29
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	155
3	ham	U dun say so early hor... U c already then say...	49
4	ham	Nah I don't think he goes to usf, he lives aro...	61

```
In [6]: mpl.rcParams['patch.force_edgecolor'] = True
        plt.style.use('seaborn-bright')
        sms.hist(column='length', by='label', bins=50, figsize=(11,5))
        plt.savefig("../img/sms-length-distribution.eps")
        plt.show()
```



Looks like the lengthy is the message, more likely it is a spam. Let’s not forget this

0.0.2 Text processing and vectorizing our meddages

Let's create new data frame. We'll need a copy later on

```
In [7]: text_feat = sms['message'].copy()
```

Now define our tex precessing function. It will remove any punctuation and stopwords aswell.

```
In [8]: def text_process(text):  
  
    text = text.translate(str.maketrans('', '', string.punctuation))  
    text = [word for word in text.split() if word.lower() not in stopwords.words('engl.  
  
    return " ".join(text)
```

```
In [9]: text_feat = text_feat.apply(text_process)
```

```
In [10]: vectorizer = TfidfVectorizer("english")
```

```
In [11]: features = vectorizer.fit_transform(text_feat)
```

```
In [12]: labels = LabelEncoder().fit_transform(sms['label'])  
labels = labels.reshape(-1,1)
```

```
In [13]: text_feat.shape
```

```
Out[13]: (5572,)
```

```
In [14]: features.shape
```

```
Out[14]: (5572, 9403)
```

0.0.3 Classifiers and predictions

First of all let's split our features to test and train set

```
In [15]: features_train, features_test, labels_train, labels_test = train_test_split(features,
```

Now let's import bunch of classifiers, initialize them and make a dictionary to itereate through

```
In [16]: from sklearn.linear_model import LogisticRegression  
from sklearn.svm import SVC  
from sklearn.naive_bayes import MultinomialNB  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.ensemble import AdaBoostClassifier  
from sklearn.ensemble import BaggingClassifier  
from sklearn.ensemble import ExtraTreesClassifier  
from sklearn.metrics import precision_score, recall_score, accuracy_score, f1_score
```

```
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/ensemble/weight_boosting.py:2
from numpy.core.umath_tests import inner1d
```

```
In [17]: svc = SVC(kernel='sigmoid', gamma=1.0)
         knc = KNeighborsClassifier(n_neighbors=49)
         mnb = MultinomialNB(alpha=0.2)
         dtc = DecisionTreeClassifier(min_samples_split=7, random_state=111)
         lrc = LogisticRegression(solver='liblinear', penalty='l1')
         rfc = RandomForestClassifier(n_estimators=31, random_state=111)
         abc = AdaBoostClassifier(n_estimators=62, random_state=111)
         bc = BaggingClassifier(n_estimators=9, random_state=111)
         etc = ExtraTreesClassifier(n_estimators=9, random_state=111)
```

```
In [18]: clfs = {'SVC' : svc, 'KN' : knc, 'NB': mnb, 'DT': dtc, 'LR': lrc, 'RF': rfc, 'AdaBoost
```

Let's make functions to fit our classifiers and make predictions

```
In [19]: def train_classifier(clf, feature_train, labels_train):
         clf.fit(feature_train, labels_train)
```

```
In [20]: def predict_labels(clf, features):
         return (clf.predict(features))
```

Now iterate through classifiers and save the results

```
In [21]: import time
```

```
In [22]: pred_scores = []
         for k,v in clfs.items():
             since = time.time()

             train_classifier(v, features_train, labels_train)
             time_elapsed = time.time() - since

             pred = predict_labels(v, features_test)
             pred_scores.append((k, [precision_score(labels_test, pred), recall_score(labels_test,
```

```
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/utils/validation.py:578: Data
    y = column_or_1d(y, warn=True)
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/ipykernel/__main__.py:2: DataConversi
    from ipykernel import kernelapp as app
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/ipykernel/__main__.py:2: DataConversi
    from ipykernel import kernelapp as app
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/utils/validation.py:578: Data
    y = column_or_1d(y, warn=True)
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/utils/validation.py:578: Data
    y = column_or_1d(y, warn=True)
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/ipykernel/__main__.py:2: DataConversi
    from ipykernel import kernelapp as app
```


```
In [23]: # pred_scores
```

```
In [24]: df = pd.DataFrame.from_items(pred_scores,orient='index', columns=['Precision', 'Recall', 'Accuracy', 'F1', 'Training Time (s)'])
df
```

```
Out[24]:
```

	Precision	Recall	Accuracy	F1	Training Time (s)
SVC	0.989305	0.825893	0.975478	0.900243	0m 0.3595s
KN	0.991304	0.508929	0.933612	0.672566	0m 0.0031s
NB	0.939535	0.901786	0.979067	0.920273	0m 0.0019s
DT	0.878378	0.870536	0.966507	0.874439	0m 0.2016s
LR	0.967532	0.665179	0.952153	0.788360	0m 0.0097s
RF	1.000000	0.781250	0.970694	0.877193	0m 1.2949s
AdaBoost	0.948980	0.830357	0.971292	0.885714	0m 2.5555s
BgC	0.905213	0.852679	0.968301	0.878161	0m 1.0332s
ETC	0.978723	0.821429	0.973684	0.893204	0m 0.9084s

```
In [25]: df.plot(kind='bar', y="Accuracy", ylim=(0.9,1.0), figsize=(11,6), align='center', col=0)
plt.xticks(np.arange(9), df.index)
plt.ylabel('Accuracy Score')
plt.title('Distribution by Classifier')
plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)
plt.savefig("../img/sms-acc-basemodel-v1-stop.eps")
plt.show()
```



Looks like ensemble classifiers are not doing as good as expected.

0.0.4 Voting classifier

We are using ensemble algorithms here, but what about ensemble of ensembles? Will it beat NB?

```

In [26]: from sklearn.ensemble import VotingClassifier

In [27]: eclf = VotingClassifier(estimators=[('BgC', bc), ('ETC', etc), ('RF', rfc), ('Ada', adaboost)])

In [28]: eclf.fit(features_train, labels_train)

/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/preprocessing/label.py:95: DataConversionWarning: Data has dtype object, but could be converted to float.
  y = column_or_1d(y, warn=True)
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/preprocessing/label.py:128: DataConversionWarning: Data has dtype object, but could be converted to float.
  y = column_or_1d(y, warn=True)

Out[28]: VotingClassifier(estimators=[('BgC', BaggingClassifier(base_estimator=None, bootstrap=True,
bootstrap_features=False, max_features=1.0, max_samples=1.0,
n_estimators=9, n_jobs=1, oob_score=False, random_state=111,
verbose=0, warm_start=False)), ('ETC', ExtraTreesClassifier(bootstrap=False,
learning_rate=1.0, n_estimators=62, random_state=111))],
flatten_transform=None, n_jobs=1, voting='soft', weights=None)

In [29]: pred = eclf.predict(features_test)

/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/preprocessing/label.py:151: DataConversionWarning: Data has dtype object, but could be converted to float.
  if diff:

In [30]: print(precision_score(labels_test, pred), recall_score(labels_test, pred), accuracy_score(labels_test, pred))

1.0 0.84375 0.979066985645933 0.9152542372881356

```

Better but nope.

0.0.5 RNN

Define the RNN structure.

```

In [31]: from keras.models import Model
         from keras.layers import LSTM, Activation, Dense, Dropout, Input, Embedding
         from keras.optimizers import RMSprop
         from keras.preprocessing.text import Tokenizer
         from keras.preprocessing import sequence
         from keras.utils import to_categorical
         from keras.callbacks import EarlyStopping
         from keras.callbacks import Callback

```

Using TensorFlow backend.

```

In [32]: max_words = features_train.shape[0]
         max_len = features_train.shape[1]

```

```
In [33]: def RNN():
    inputs = Input(name='inputs',shape=[max_len])
    layer = Embedding(max_words,50,input_length=max_len)(inputs)
    layer = LSTM(64)(layer)
    layer = Dense(256,name='FC1')(layer)
    layer = Activation('relu')(layer)
    layer = Dropout(0.5)(layer)
    layer = Dense(1,name='out_layer')(layer)
    layer = Activation('sigmoid')(layer)
    model = Model(inputs=inputs,outputs=layer)
    return model
```

Call the function and compile the model.

```
In [34]: model = RNN()
    model.summary()
    model.compile(loss='binary_crossentropy',optimizer=RMSprop(),metrics=['accuracy'])
```

```
-----
Layer (type)                 Output Shape              Param #
=====
inputs (InputLayer)          (None, 9403)              0
-----
embedding_1 (Embedding)      (None, 9403, 50)         195000
-----
lstm_1 (LSTM)                 (None, 64)                29440
-----
FC1 (Dense)                   (None, 256)               16640
-----
activation_1 (Activation)     (None, 256)               0
-----
dropout_1 (Dropout)           (None, 256)               0
-----
out_layer (Dense)             (None, 1)                 257
-----
activation_2 (Activation)     (None, 1)                 0
=====
Total params: 241,337
Trainable params: 241,337
Non-trainable params: 0
-----
```

```
In [35]: since = time.time()

    # model.fit(features_train, labels_train, epochs=10, batch_size=128,
    #           validation_split=0.2,
    #           callbacks=[metrics, EarlyStopping(monitor='val_loss',min_delt
```

```

model.fit(features_train, labels_train, batch_size=128, epochs=10,
          validation_split=0.2, callbacks=[EarlyStopping(monitor='val_loss', min_delta=

time_elapsed = time.time() - since

Train on 3120 samples, validate on 780 samples
Epoch 1/10
3120/3120 [=====] - 298s 96ms/step - loss: 0.4477 - acc: 0.8471 - val_
Epoch 2/10
3120/3120 [=====] - 290s 93ms/step - loss: 0.4053 - acc: 0.8654 - val_

In [36]: print('Training complete in {:.0f}m {:.4f}s'.format(
          time_elapsed // 60, time_elapsed % 60))

```

Training complete in 9m 49.5561s

```

In [38]: pred = (np.asarray(model.predict(features_test, batch_size=128))).round()

In [42]: pred_scores.append(("LSTM", [precision_score(labels_test, pred), recall_score(labels_test, pred)]))

/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/metrics/classification.py:113:
  'precision', 'predicted', average, warn_for)
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/metrics/classification.py:113:
  'precision', 'predicted', average, warn_for)

```

0.0.6 gcForest

```

In [43]: import sys
          sys.path.append("..")
          from gcforest.gcforest import GCForest
          from gcforest.utils.config_utils import load_json

In [44]: def get_toy_config():
          config = {}
          ca_config = {}
          ca_config["random_state"] = 111
          ca_config["max_layers"] = 10
          ca_config["early_stopping_rounds"] = 3
          ca_config["n_classes"] = 2
          ca_config["estimators"] = []
          ca_config["estimators"].append({"n_folds": 5, "type": "RandomForestClassifier", "n_estimators": 100})
          ca_config["estimators"].append({"n_folds": 5, "type": "RandomForestClassifier", "n_estimators": 100})
          ca_config["estimators"].append({"n_folds": 5, "type": "MultinomialNB", "alpha": 0.01})
          ca_config["estimators"].append({"n_folds": 5, "type": "MultinomialNB", "alpha": 0.01})

          config["cascade"] = ca_config
          return config

```



```

In [45]: config = get_toy_config()
         gc = GCForest(config)

         # features_train ndarraylabels_train (n_samples, )(n_samples, 1)
         features_train = features_train.toarray()
         labels_train = labels_train.reshape(-1)

         since = time.time()
         gc.fit_transform(features_train, labels_train)

         time_elapsed = time.time() - since

[ 2019-04-23 21:00:02,205] [cascade_classifier.fit_transform] X_groups_train.shape=[(3900, 9403)
[ 2019-04-23 21:00:02,455] [cascade_classifier.fit_transform] group_dims=[9403]
[ 2019-04-23 21:00:02,456] [cascade_classifier.fit_transform] group_starts=[0]
[ 2019-04-23 21:00:02,456] [cascade_classifier.fit_transform] group_ends=[9403]
[ 2019-04-23 21:00:02,457] [cascade_classifier.fit_transform] X_train.shape=(3900, 9403),X_test
[ 2019-04-23 21:00:02,645] [cascade_classifier.fit_transform] [layer=0] look_indexes=[0], X_cur_t
[ 2019-04-23 21:00:06,792] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_0 - 5_
[ 2019-04-23 21:00:10,442] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_0 - 5_
[ 2019-04-23 21:00:13,723] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_0 - 5_
[ 2019-04-23 21:00:17,217] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_0 - 5_
[ 2019-04-23 21:00:20,761] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_0 - 5_
[ 2019-04-23 21:00:20,762] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_0 - 5_
[ 2019-04-23 21:00:24,097] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_1 - 5_
[ 2019-04-23 21:00:27,550] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_1 - 5_
[ 2019-04-23 21:00:30,885] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_1 - 5_
[ 2019-04-23 21:00:34,334] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_1 - 5_
[ 2019-04-23 21:00:37,672] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_1 - 5_
[ 2019-04-23 21:00:37,673] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_1 - 5_
[ 2019-04-23 21:00:37,909] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_2 - 5_
[ 2019-04-23 21:00:38,141] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_2 - 5_
[ 2019-04-23 21:00:38,367] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_2 - 5_
[ 2019-04-23 21:00:38,593] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_2 - 5_
[ 2019-04-23 21:00:38,822] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_2 - 5_
[ 2019-04-23 21:00:38,823] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_2 - 5_
[ 2019-04-23 21:00:39,049] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_3 - 5_
[ 2019-04-23 21:00:39,276] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_3 - 5_
[ 2019-04-23 21:00:39,501] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_3 - 5_
[ 2019-04-23 21:00:39,722] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_3 - 5_
[ 2019-04-23 21:00:39,946] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_3 - 5_
[ 2019-04-23 21:00:39,947] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_3 - 5_
[ 2019-04-23 21:00:39,948] [cascade_classifier.calc_accuracy] Accuracy(layer_0 - train.classifi
[ 2019-04-23 21:00:40,154] [cascade_classifier.fit_transform] [layer=1] look_indexes=[0], X_cur_t
[ 2019-04-23 21:00:42,324] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_0 - 5_
[ 2019-04-23 21:00:44,332] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_0 - 5_
[ 2019-04-23 21:00:46,486] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_0 - 5_
[ 2019-04-23 21:00:48,549] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_0 - 5_

```

```

[ 2019-04-23 21:00:50,349] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_0 - 5_
[ 2019-04-23 21:00:50,350] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_0 - 5_
[ 2019-04-23 21:00:52,456] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_1 - 5_
[ 2019-04-23 21:00:54,561] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_1 - 5_
[ 2019-04-23 21:00:56,455] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_1 - 5_
[ 2019-04-23 21:00:58,472] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_1 - 5_
[ 2019-04-23 21:01:00,569] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_1 - 5_
[ 2019-04-23 21:01:00,571] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_1 - 5_
[ 2019-04-23 21:01:00,810] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_2 - 5_
[ 2019-04-23 21:01:01,038] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_2 - 5_
[ 2019-04-23 21:01:01,267] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_2 - 5_
[ 2019-04-23 21:01:01,491] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_2 - 5_
[ 2019-04-23 21:01:01,722] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_2 - 5_
[ 2019-04-23 21:01:01,723] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_2 - 5_
[ 2019-04-23 21:01:01,965] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_3 - 5_
[ 2019-04-23 21:01:02,191] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_3 - 5_
[ 2019-04-23 21:01:02,419] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_3 - 5_
[ 2019-04-23 21:01:02,644] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_3 - 5_
[ 2019-04-23 21:01:02,876] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_3 - 5_
[ 2019-04-23 21:01:02,877] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_3 - 5_
[ 2019-04-23 21:01:02,878] [cascade_classifier.calc_accuracy] Accuracy(layer_1 - train.classifi
[ 2019-04-23 21:01:03,059] [cascade_classifier.fit_transform] [layer=2] look_indexes=[0], X_cur_t
[ 2019-04-23 21:01:04,857] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_0 - 5_
[ 2019-04-23 21:01:06,762] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_0 - 5_
[ 2019-04-23 21:01:08,654] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_0 - 5_
[ 2019-04-23 21:01:10,437] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_0 - 5_
[ 2019-04-23 21:01:12,226] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_0 - 5_
[ 2019-04-23 21:01:12,227] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_0 - 5_
[ 2019-04-23 21:01:14,023] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_1 - 5_
[ 2019-04-23 21:01:15,795] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_1 - 5_
[ 2019-04-23 21:01:17,683] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_1 - 5_
[ 2019-04-23 21:01:19,476] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_1 - 5_
[ 2019-04-23 21:01:21,268] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_1 - 5_
[ 2019-04-23 21:01:21,270] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_1 - 5_
[ 2019-04-23 21:01:21,502] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_2 - 5_
[ 2019-04-23 21:01:21,729] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_2 - 5_
[ 2019-04-23 21:01:22,070] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_2 - 5_
[ 2019-04-23 21:01:22,292] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_2 - 5_
[ 2019-04-23 21:01:22,514] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_2 - 5_
[ 2019-04-23 21:01:22,515] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_2 - 5_
[ 2019-04-23 21:01:22,741] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_3 - 5_
[ 2019-04-23 21:01:22,960] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_3 - 5_
[ 2019-04-23 21:01:23,183] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_3 - 5_
[ 2019-04-23 21:01:23,509] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_3 - 5_
[ 2019-04-23 21:01:23,731] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_3 - 5_
[ 2019-04-23 21:01:23,732] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_3 - 5_
[ 2019-04-23 21:01:23,733] [cascade_classifier.calc_accuracy] Accuracy(layer_2 - train.classifi
[ 2019-04-23 21:01:23,919] [cascade_classifier.fit_transform] [layer=3] look_indexes=[0], X_cur_t

```

```

[ 2019-04-23 21:01:25,814] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_0 - 5_1
[ 2019-04-23 21:01:27,814] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_0 - 5_1
[ 2019-04-23 21:01:29,804] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_0 - 5_1
[ 2019-04-23 21:01:31,581] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_0 - 5_1
[ 2019-04-23 21:01:33,575] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_0 - 5_1
[ 2019-04-23 21:01:33,579] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_0 - 5_1
[ 2019-04-23 21:01:35,461] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_1 - 5_1
[ 2019-04-23 21:01:37,341] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_1 - 5_1
[ 2019-04-23 21:01:39,229] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_1 - 5_1
[ 2019-04-23 21:01:41,114] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_1 - 5_1
[ 2019-04-23 21:01:43,103] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_1 - 5_1
[ 2019-04-23 21:01:43,104] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_1 - 5_1
[ 2019-04-23 21:01:43,330] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_2 - 5_1
[ 2019-04-23 21:01:43,552] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_2 - 5_1
[ 2019-04-23 21:01:43,776] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_2 - 5_1
[ 2019-04-23 21:01:43,998] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_2 - 5_1
[ 2019-04-23 21:01:44,221] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_2 - 5_1
[ 2019-04-23 21:01:44,222] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_2 - 5_1
[ 2019-04-23 21:01:44,442] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_3 - 5_1
[ 2019-04-23 21:01:44,663] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_3 - 5_1
[ 2019-04-23 21:01:44,883] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_3 - 5_1
[ 2019-04-23 21:01:45,103] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_3 - 5_1
[ 2019-04-23 21:01:45,322] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_3 - 5_1
[ 2019-04-23 21:01:45,323] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_3 - 5_1
[ 2019-04-23 21:01:45,325] [cascade_classifier.calc_accuracy] Accuracy(layer_3 - train.classific
[ 2019-04-23 21:01:45,509] [cascade_classifier.fit_transform] [layer=4] look_indexes=[0], X_cur_t
[ 2019-04-23 21:01:47,608] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_0 - 5_1
[ 2019-04-23 21:01:49,600] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_0 - 5_1
[ 2019-04-23 21:01:51,584] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_0 - 5_1
[ 2019-04-23 21:01:53,370] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_0 - 5_1
[ 2019-04-23 21:01:55,457] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_0 - 5_1
[ 2019-04-23 21:01:55,458] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_0 - 5_1
[ 2019-04-23 21:01:57,349] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_1 - 5_1
[ 2019-04-23 21:01:59,229] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_1 - 5_1
[ 2019-04-23 21:02:01,229] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_1 - 5_1
[ 2019-04-23 21:02:03,319] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_1 - 5_1
[ 2019-04-23 21:02:05,304] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_1 - 5_1
[ 2019-04-23 21:02:05,306] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_1 - 5_1
[ 2019-04-23 21:02:05,525] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_2 - 5_1
[ 2019-04-23 21:02:05,744] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_2 - 5_1
[ 2019-04-23 21:02:05,966] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_2 - 5_1
[ 2019-04-23 21:02:06,189] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_2 - 5_1
[ 2019-04-23 21:02:06,409] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_2 - 5_1
[ 2019-04-23 21:02:06,410] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_2 - 5_1
[ 2019-04-23 21:02:06,635] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_3 - 5_1
[ 2019-04-23 21:02:06,857] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_3 - 5_1
[ 2019-04-23 21:02:07,073] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_3 - 5_1
[ 2019-04-23 21:02:07,294] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_3 - 5_1

```

```
[ 2019-04-23 21:02:07,511][kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_3 - 5_1
[ 2019-04-23 21:02:07,512][kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_3 - 5_1
[ 2019-04-23 21:02:07,514][cascade_classifier.calc_accuracy] Accuracy(layer_4 - train.classifi
[ 2019-04-23 21:02:07,515][cascade_classifier.fit_transform] [Result][Optimal Level Detected] c
```

```
In [46]: print('Training complete in {:.0f}m {:.4f}s'.format(
            time_elapsed // 60, time_elapsed % 60))
```

Training complete in 2m 5.3713s

```
In [47]: pred = predict_labels(gc,features_test.toarray())
        pred_scores.append(("DCF", [precision_score(labels_test,pred), recall_score(labels_test,
```


```
[ 2019-04-23 21:02:07,611][cascade_classifier.transform] X_groups_test.shape=[[1672, 9403]]
[ 2019-04-23 21:02:07,740][cascade_classifier.transform] group_dims=[9403]
[ 2019-04-23 21:02:07,740][cascade_classifier.transform] X_test.shape=(1672, 9403)
[ 2019-04-23 21:02:07,818][cascade_classifier.transform] [layer=0] look_indexes=[0], X_cur_test
[ 2019-04-23 21:02:09,504][cascade_classifier.transform] [layer=1] look_indexes=[0], X_cur_test
```

```
In [48]: df = pd.DataFrame.from_items(pred_scores,orient='index', columns=['Precision', 'Recall',
        df
```

```
Out[48]:
```

	Precision	Recall	Accuracy	F1	Training Time (s)
SVC	0.989305	0.825893	0.975478	0.900243	0m 0.3595s
KN	0.991304	0.508929	0.933612	0.672566	0m 0.0031s
NB	0.939535	0.901786	0.979067	0.920273	0m 0.0019s
DT	0.878378	0.870536	0.966507	0.874439	0m 0.2016s
LR	0.967532	0.665179	0.952153	0.788360	0m 0.0097s
RF	1.000000	0.781250	0.970694	0.877193	0m 1.2949s
AdaBoost	0.948980	0.830357	0.971292	0.885714	0m 2.5555s
BgC	0.905213	0.852679	0.968301	0.878161	0m 1.0332s
ETC	0.978723	0.821429	0.973684	0.893204	0m 0.9084s
LSTM	0.000000	0.000000	0.866029	0.000000	9m 49.5561s
DCF	0.980952	0.919643	0.986842	0.949309	2m 5.3713s

```
In [49]: df.plot(kind='bar', y="Accuracy", ylim=(0.8,1.0), figsize=(11,6), align='center', col
        plt.xticks(np.arange(11), df.index)
        plt.ylabel('Accuracy Score')
        plt.title('Distribution by Classifier')
        plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)
        plt.savefig("../img/sms-acc-v1-stop.eps")
        plt.show()
```



```
In [50]: import pickle
# dump
with open("../pkl/sms-gc-v1-stop.pkl", "wb") as f:
    pickle.dump(gc, f, pickle.HIGHEST_PROTOCOL)

# # load
# with open("../pkl/2018_gc.pkl", "rb") as f:
#     gc = pickle.load(f)
```

0.0.7 Final verdict - gcForest is your friend in spam detection.

```
In [ ]:
```