

sms-spam-v3

April 23, 2019

V3:

- LSTM uses the `Tokenizer.fit_on_texts(data)` then `Tokenizer.texts_to_sequences(data)` to do the preprocessing
- Other models uses the `TfidfVectorizer` to do the the preprocessing

Goal of this notebook to test several classifiers on the data set with different features

0.0.1 Let's begin

First of all necessary imports

```
In [1]: import numpy as np
import pandas as pd
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sns
import string
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from nltk.corpus import stopwords
from sklearn.preprocessing import LabelEncoder
%matplotlib inline
```

Let's read the data from csv file

```
In [2]: sms = pd.read_csv('../data/sms-spam.csv', delimiter=',', encoding='latin-1')

sms.head()
```

```
Out[2]:
```

	v1	v2	Unnamed: 2	\
0	ham	Go until jurong point, crazy.. Available only ...	NaN	
1	ham	Ok lar... Joking wif u oni...	NaN	
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	
3	ham	U dun say so early hor... U c already then say...	NaN	
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	

Unnamed: 3 Unnamed: 4

0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

Now drop “unnamed” columns and rename v1 and v2 to “label” and “message”

```
In [3]: sms = sms.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], axis=1)
        sms = sms.rename(columns = {'v1': 'label', 'v2': 'message'})
```

Let’s look into our data

```
In [4]: sms.groupby('label').describe()
```

```
Out[4]:
```

	message								
		count	unique					top	freq
label									
ham		4825	4516				Sorry, I'll call later	30	
spam		747	653	Please call our customer service representativ...				4	

Intresting that “Sorry, I’ll call later” appears only 30 times here =)


Now let’s create new feature “message length” and plot it to see if it’s of any interest

```
In [5]: sms['length'] = sms['message'].apply(len)
        sms.head()
```

```
Out[5]:
```

	label	message	length
0	ham	Go until jurong point, crazy.. Available only ...	111
1	ham	Ok lar... Joking wif u oni...	29
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	155
3	ham	U dun say so early hor... U c already then say...	49
4	ham	Nah I don't think he goes to usf, he lives aro...	61

```
In [6]: mpl.rcParams['patch.force_edgecolor'] = True
        plt.style.use('seaborn-bright')
        sms.hist(column='length', by='label', bins=50, figsize=(11,5))
        plt.savefig("../img/sms-length-distribution.eps")
        plt.show()
```



Looks like the lengthy is the message, more likely it is a spam. Let's not forget this

0.0.2 Text processing and vectorizing our meddages

Let's create new data frame. We'll need a copy later on

```
In [7]: text_feat = sms['message'].copy()
```

Now define our tex precessing function. It will remove any punctuation and stopwords aswell.

```
In [8]: # def text_process(text):

#     text = text.translate(str.maketrans('', '', string.punctuation))
#     text = [word for word in text.split() if word.lower() not in stopwords.words('en')]

#     return " ".join(text)
```

```
In [9]: # text_feat = text_feat.apply(text_process)
```

```
In [10]: vectorizer = TfidfVectorizer("english")
```

```
In [11]: features = vectorizer.fit_transform(text_feat)
```

```
In [12]: labels = LabelEncoder().fit_transform(sms['label'])
labels = labels.reshape(-1,1)
```

```
In [13]: text_feat.shape
```

```
Out[13]: (5572,)
```

```
In [14]: features.shape
```

```
Out[14]: (5572, 8710)
```

0.0.3 Classifiers and predictions

First of all let's split our features to test and train set

```
In [15]: features_train, features_test, labels_train, labels_test = train_test_split(features,
```

Now let's import bunch of classifiers, initialize them and make a dictionary to iterate through

```
In [16]: from sklearn.linear_model import LogisticRegression
         from sklearn.svm import SVC
         from sklearn.naive_bayes import MultinomialNB
         from sklearn.tree import DecisionTreeClassifier
         from sklearn.neighbors import KNeighborsClassifier
         from sklearn.ensemble import RandomForestClassifier
         from sklearn.ensemble import AdaBoostClassifier
         from sklearn.ensemble import BaggingClassifier
         from sklearn.ensemble import ExtraTreesClassifier
         from sklearn.metrics import precision_score, recall_score, accuracy_score, f1_score
```

```
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/ensemble/weight_boosting.py:2
         from numpy.core.umath_tests import inner1d
```

```
In [17]: svc = SVC(kernel='sigmoid', gamma=1.0)
         knc = KNeighborsClassifier(n_neighbors=49)
         mnb = MultinomialNB(alpha=0.2)
         dtc = DecisionTreeClassifier(min_samples_split=7, random_state=111)
         lrc = LogisticRegression(solver='liblinear', penalty='l1')
         rfc = RandomForestClassifier(n_estimators=31, random_state=111)
         abc = AdaBoostClassifier(n_estimators=62, random_state=111)
         bc = BaggingClassifier(n_estimators=9, random_state=111)
         etc = ExtraTreesClassifier(n_estimators=9, random_state=111)
```

```
In [18]: clfs = {'SVC' : svc, 'KN' : knc, 'NB': mnb, 'DT': dtc, 'LR': lrc, 'RF': rfc, 'AdaBoost
```

Let's make functions to fit our classifiers and make predictions

```
In [19]: def train_classifier(clf, feature_train, labels_train):
         clf.fit(feature_train, labels_train)
```

```
In [20]: def predict_labels(clf, features):
         return (clf.predict(features))
```

Now iterate through classifiers and save the results

```
In [21]: import time
```

```
In [22]: pred_scores = []
         for k,v in clfs.items():
             since = time.time()
```

```

train_classifier(v, features_train, labels_train)
time_elapsed = time.time() - since

pred = predict_labels(v, features_test)
pred_scores.append((k, [precision_score(labels_test, pred), recall_score(labels_test, pred)]))

/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/utils/validation.py:578: DataConversionWarning: A column-vector y was passed when you expected a scalar:
y = column_or_1d(y, warn=True)
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/ipykernel/__main__.py:2: DataConversionWarning: A column-vector y was passed when you expected a scalar:
from ipykernel import kernelapp as app
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/ipykernel/__main__.py:2: DataConversionWarning: A column-vector y was passed when you expected a scalar:
from ipykernel import kernelapp as app
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/utils/validation.py:578: DataConversionWarning: A column-vector y was passed when you expected a scalar:
y = column_or_1d(y, warn=True)
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/utils/validation.py:578: DataConversionWarning: A column-vector y was passed when you expected a scalar:
y = column_or_1d(y, warn=True)
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/ipykernel/__main__.py:2: DataConversionWarning: A column-vector y was passed when you expected a scalar:
from ipykernel import kernelapp as app

```

In [23]: # *pred_scores*

```
In [24]: df = pd.DataFrame.from_items(pred_scores, orient='index', columns=['Precision', 'Recall', 'Accuracy', 'F1', 'Training Time (s)'])
df
```


```
Out[24]:
```

	Precision	Recall	Accuracy	F1	Training Time (s)
SVC	0.990476	0.896552	0.984450	0.941176	0m 0.4593s
KN	1.000000	0.616379	0.946770	0.762667	0m 0.0041s
NB	0.982143	0.948276	0.990431	0.964912	0m 0.0018s
DT	0.900452	0.857759	0.967105	0.878587	0m 0.2270s
LR	0.933333	0.724138	0.954545	0.815534	0m 0.0107s
RF	1.000000	0.866379	0.981459	0.928406	0m 0.9180s
AdaBoost	0.952607	0.866379	0.975478	0.907449	0m 2.4091s
BgC	0.935185	0.870690	0.973684	0.901786	0m 1.1530s
ETC	0.995025	0.862069	0.980263	0.923788	0m 0.6672s

```

In [25]: df.plot(kind='bar', y="Accuracy", ylim=(0.9,1.0), figsize=(11,6), align='center', color='red')
plt.xticks(np.arange(9), df.index)
plt.ylabel('Accuracy Score')
plt.title('Distribution by Classifier')
plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)
plt.savefig("../img/sms-acc-basemodel-v3.eps")
plt.show()

```



Looks like ensemble classifiers are not doing as good as expected.

0.0.4 Voting classifier

We are using ensemble algorithms here, but what about ensemble of ensembles? Will it beat NB?

```
In [26]: from sklearn.ensemble import VotingClassifier
```

```
In [27]: eclf = VotingClassifier(estimators=[('BgC', bc), ('ETC', etc), ('RF', rfc), ('Ada', al)
```

```
In [28]: eclf.fit(features_train, labels_train)
```

```
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/preprocessing/label.py:95: DataConversionWarning:
  y = column_or_1d(y, warn=True)
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/preprocessing/label.py:128: DataConversionWarning:
  y = column_or_1d(y, warn=True)
```

```
Out[28]: VotingClassifier(estimators=[('BgC', BaggingClassifier(base_estimator=None, bootstrap=
  bootstrap_features=False, max_features=1.0, max_samples=1.0,
  n_estimators=9, n_jobs=1, oob_score=False, random_state=111,
  verbose=0, warm_start=False)), ('ETC', ExtraTreesClassifier(bootstrap=False,
  learning_rate=1.0, n_estimators=62, random_state=111))],
  flatten_transform=None, n_jobs=1, voting='soft', weights=None)
```

```
In [29]: pred = eclf.predict(features_test)
```

```
/Users/alex/anaconda/envs/gc/lib/python3.6/site-packages/sklearn/preprocessing/label.py:151: DataConversionWarning:
  if diff:
```

```
In [30]: print(precision_score(labels_test,pred), recall_score(labels_test,pred), accuracy_score(labels_test,pred))
0.9806763285024155 0.875 0.9802631578947368 0.9248291571753987
```

Better but nope.

0.0.5 RNN

Define the RNN structure.

```
In [31]: from keras.models import Model
        from keras.layers import LSTM, Activation, Dense, Dropout, Input, Embedding
        from keras.optimizers import RMSprop
        from keras.preprocessing.text import Tokenizer
        from keras.preprocessing import sequence
        from keras.utils import to_categorical
        from keras.callbacks import EarlyStopping
        from keras.callbacks import Callback
```

Using TensorFlow backend.

0.0.6 Process the data

- Tokenize the data and convert the text to sequences.
- Add padding to ensure that all the sequences have the same shape.
- There are many ways of taking the *max_len* and here an arbitrary length of 500 is chosen. (From the Fig, almost all the sentences have the length < 200)

```
In [32]: features_lstm = text_feat
        labels_lstm = labels
```

```
In [33]: max_words = 1000
        max_len = 200 # n_features
        tok = Tokenizer(num_words=max_words)
        tok.fit_on_texts(features_lstm)
        sequences = tok.texts_to_sequences(features_lstm)
        features_lstm = sequence.pad_sequences(sequences,maxlen=max_len)
```

```
In [34]: features_lstm.shape
```

```
Out[34]: (5572, 200)
```

```
In [35]: labels_lstm.shape
```

```
Out[35]: (5572, 1)
```

```
In [36]: features_lstm_train, features_lstm_test, labels_lstm_train, labels_lstm_test = train_test_split(features_lstm, labels_lstm, test_size=0.1, random_state=42)
```

```
In [37]: def RNN():
    inputs = Input(name='inputs',shape=[max_len])
    layer = Embedding(max_words,50,input_length=max_len)(inputs)
    layer = LSTM(64)(layer)
    layer = Dense(256,name='FC1')(layer)
    layer = Activation('relu')(layer)
    layer = Dropout(0.5)(layer)
    layer = Dense(1,name='out_layer')(layer)
    layer = Activation('sigmoid')(layer)
    model = Model(inputs=inputs,outputs=layer)
    return model
```

Call the function and compile the model.

```
In [38]: model = RNN()
    model.summary()
    model.compile(loss='binary_crossentropy',optimizer=RMSprop(),metrics=['accuracy'])
```

Layer (type)	Output Shape	Param #
inputs (InputLayer)	(None, 200)	0
embedding_1 (Embedding)	(None, 200, 50)	50000
lstm_1 (LSTM)	(None, 64)	29440
FC1 (Dense)	(None, 256)	16640
activation_1 (Activation)	(None, 256)	0
dropout_1 (Dropout)	(None, 256)	0
out_layer (Dense)	(None, 1)	257
activation_2 (Activation)	(None, 1)	0

Total params: 96,337
 Trainable params: 96,337
 Non-trainable params: 0

```
In [39]: since = time.time()
```

```
    model.fit(features_lstm_train, labels_lstm_train, epochs=10, batch_size=128, validation_data=(features_lstm_val, labels_lstm_val),
              callbacks=[EarlyStopping(monitor='val_loss',min_delta=0.0001)])
```



```

        time_elapsed = time.time() - since

Train on 3120 samples, validate on 780 samples
Epoch 1/10
3120/3120 [=====] - 9s 3ms/step - loss: 0.3671 - acc: 0.8538 - val_loss: 0.1145
Epoch 2/10
3120/3120 [=====] - 8s 3ms/step - loss: 0.1145 - acc: 0.9696 - val_loss: 0.0458
Epoch 3/10
3120/3120 [=====] - 6s 2ms/step - loss: 0.0458 - acc: 0.9869 - val_loss: 0.0114

```

```

In [40]: print('Training complete in {:.0f}m {:.4f}s'.format(
            time_elapsed // 60, time_elapsed % 60))

```

Training complete in 0m 23.4645s

```

In [41]: pred = (np.asarray(model.predict(features_lstm_test, batch_size=128))).round()

```

```

In [42]: pred_scores.append(("LSTM", [precision_score(labels_lstm_test, pred), recall_score(labels_lstm_test, pred)]))

```

0.0.7 gcForest

```

In [43]: import sys
         sys.path.append("..")
         from gcforest.gcforest import GCForest
         from gcforest.utils.config_utils import load_json

In [44]: def get_toy_config():
         config = {}
         ca_config = {}
         ca_config["random_state"] = 111
         ca_config["max_layers"] = 10
         ca_config["early_stopping_rounds"] = 3
         ca_config["n_classes"] = 2
         ca_config["estimators"] = []
         ca_config["estimators"].append({"n_folds": 5, "type": "RandomForestClassifier", "n_estimators": 100})
         ca_config["estimators"].append({"n_folds": 5, "type": "MultinomialNB", "alpha": 0.01})

         config["cascade"] = ca_config
         return config

In [45]: config = get_toy_config()
         gc = GCForest(config)

         # features_train ndarraylabels_train (n_samples, )(n_samples, 1)
         features_train = features_train.toarray()
         labels_train = labels_train.reshape(-1)

```

```

since = time.time()
gc.fit_transform(features_train, labels_train)

time_elapsed = time.time() - since

# gc.fit_transform(features_train, labels_train, features_test, labels_test)

[ 2019-04-23 22:00:22,133] [cascade_classifier.fit_transform] X_groups_train.shape=[(3900, 8710)
[ 2019-04-23 22:00:22,356] [cascade_classifier.fit_transform] group_dims=[8710]
[ 2019-04-23 22:00:22,358] [cascade_classifier.fit_transform] group_starts=[0]
[ 2019-04-23 22:00:22,358] [cascade_classifier.fit_transform] group_ends=[8710]
[ 2019-04-23 22:00:22,359] [cascade_classifier.fit_transform] X_train.shape=(3900, 8710),X_test
[ 2019-04-23 22:00:22,527] [cascade_classifier.fit_transform] [layer=0] look_indexs=[0], X_cur
[ 2019-04-23 22:00:25,099] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_0 - 5_
[ 2019-04-23 22:00:27,622] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_0 - 5_
[ 2019-04-23 22:00:30,116] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_0 - 5_
[ 2019-04-23 22:00:32,607] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_0 - 5_
[ 2019-04-23 22:00:35,006] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_0 - 5_
[ 2019-04-23 22:00:35,007] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_0 - 5_
[ 2019-04-23 22:00:35,218] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_1 - 5_
[ 2019-04-23 22:00:35,427] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_1 - 5_
[ 2019-04-23 22:00:35,634] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_1 - 5_
[ 2019-04-23 22:00:35,844] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_1 - 5_
[ 2019-04-23 22:00:36,051] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_1 - 5_
[ 2019-04-23 22:00:36,052] [kfold_wrapper.log_eval_metrics] Accuracy(layer_0 - estimator_1 - 5_
[ 2019-04-23 22:00:36,053] [cascade_classifier.calc_accuracy] Accuracy(layer_0 - train.classifi
[ 2019-04-23 22:00:36,227] [cascade_classifier.fit_transform] [layer=1] look_indexs=[0], X_cur
[ 2019-04-23 22:00:38,326] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_0 - 5_
[ 2019-04-23 22:00:40,314] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_0 - 5_
[ 2019-04-23 22:00:42,200] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_0 - 5_
[ 2019-04-23 22:00:44,191] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_0 - 5_
[ 2019-04-23 22:00:45,962] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_0 - 5_
[ 2019-04-23 22:00:45,964] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_0 - 5_
[ 2019-04-23 22:00:46,181] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_1 - 5_
[ 2019-04-23 22:00:46,397] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_1 - 5_
[ 2019-04-23 22:00:46,614] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_1 - 5_
[ 2019-04-23 22:00:46,832] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_1 - 5_
[ 2019-04-23 22:00:47,049] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_1 - 5_
[ 2019-04-23 22:00:47,050] [kfold_wrapper.log_eval_metrics] Accuracy(layer_1 - estimator_1 - 5_
[ 2019-04-23 22:00:47,052] [cascade_classifier.calc_accuracy] Accuracy(layer_1 - train.classifi
[ 2019-04-23 22:00:47,240] [cascade_classifier.fit_transform] [layer=2] look_indexs=[0], X_cur
[ 2019-04-23 22:00:48,928] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_0 - 5_
[ 2019-04-23 22:00:50,611] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_0 - 5_
[ 2019-04-23 22:00:52,392] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_0 - 5_
[ 2019-04-23 22:00:54,081] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_0 - 5_
[ 2019-04-23 22:00:55,766] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_0 - 5_
[ 2019-04-23 22:00:55,767] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_0 - 5_

```

```

[ 2019-04-23 22:00:55,988] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_1 - 5_)
[ 2019-04-23 22:00:56,206] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_1 - 5_)
[ 2019-04-23 22:00:56,424] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_1 - 5_)
[ 2019-04-23 22:00:56,642] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_1 - 5_)
[ 2019-04-23 22:00:56,859] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_1 - 5_)
[ 2019-04-23 22:00:56,860] [kfold_wrapper.log_eval_metrics] Accuracy(layer_2 - estimator_1 - 5_)
[ 2019-04-23 22:00:56,861] [cascade_classifier.calc_accuracy] Accuracy(layer_2 - train.classifi
[ 2019-04-23 22:00:57,031] [cascade_classifier.fit_transform] [layer=3] look_indexs=[0], X_cur
[ 2019-04-23 22:00:58,823] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_0 - 5_)
[ 2019-04-23 22:01:00,703] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_0 - 5_)
[ 2019-04-23 22:01:02,374] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_0 - 5_)
[ 2019-04-23 22:01:04,166] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_0 - 5_)
[ 2019-04-23 22:01:05,834] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_0 - 5_)
[ 2019-04-23 22:01:05,835] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_0 - 5_)
[ 2019-04-23 22:01:06,055] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_1 - 5_)
[ 2019-04-23 22:01:06,276] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_1 - 5_)
[ 2019-04-23 22:01:06,494] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_1 - 5_)
[ 2019-04-23 22:01:06,719] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_1 - 5_)
[ 2019-04-23 22:01:06,935] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_1 - 5_)
[ 2019-04-23 22:01:06,936] [kfold_wrapper.log_eval_metrics] Accuracy(layer_3 - estimator_1 - 5_)
[ 2019-04-23 22:01:06,937] [cascade_classifier.calc_accuracy] Accuracy(layer_3 - train.classifi
[ 2019-04-23 22:01:07,166] [cascade_classifier.fit_transform] [layer=4] look_indexs=[0], X_cur
[ 2019-04-23 22:01:08,761] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_0 - 5_)
[ 2019-04-23 22:01:10,336] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_0 - 5_)
[ 2019-04-23 22:01:12,021] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_0 - 5_)
[ 2019-04-23 22:01:13,590] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_0 - 5_)
[ 2019-04-23 22:01:15,261] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_0 - 5_)
[ 2019-04-23 22:01:15,262] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_0 - 5_)
[ 2019-04-23 22:01:15,472] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_1 - 5_)
[ 2019-04-23 22:01:15,680] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_1 - 5_)
[ 2019-04-23 22:01:15,893] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_1 - 5_)
[ 2019-04-23 22:01:16,103] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_1 - 5_)
[ 2019-04-23 22:01:16,314] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_1 - 5_)
[ 2019-04-23 22:01:16,315] [kfold_wrapper.log_eval_metrics] Accuracy(layer_4 - estimator_1 - 5_)
[ 2019-04-23 22:01:16,316] [cascade_classifier.calc_accuracy] Accuracy(layer_4 - train.classifi
[ 2019-04-23 22:01:16,489] [cascade_classifier.fit_transform] [layer=5] look_indexs=[0], X_cur
[ 2019-04-23 22:01:18,274] [kfold_wrapper.log_eval_metrics] Accuracy(layer_5 - estimator_0 - 5_)
[ 2019-04-23 22:01:19,942] [kfold_wrapper.log_eval_metrics] Accuracy(layer_5 - estimator_0 - 5_)
[ 2019-04-23 22:01:21,401] [kfold_wrapper.log_eval_metrics] Accuracy(layer_5 - estimator_0 - 5_)
[ 2019-04-23 22:01:22,966] [kfold_wrapper.log_eval_metrics] Accuracy(layer_5 - estimator_0 - 5_)
[ 2019-04-23 22:01:24,641] [kfold_wrapper.log_eval_metrics] Accuracy(layer_5 - estimator_0 - 5_)
[ 2019-04-23 22:01:24,643] [kfold_wrapper.log_eval_metrics] Accuracy(layer_5 - estimator_0 - 5_)
[ 2019-04-23 22:01:24,847] [kfold_wrapper.log_eval_metrics] Accuracy(layer_5 - estimator_1 - 5_)
[ 2019-04-23 22:01:25,052] [kfold_wrapper.log_eval_metrics] Accuracy(layer_5 - estimator_1 - 5_)
[ 2019-04-23 22:01:25,259] [kfold_wrapper.log_eval_metrics] Accuracy(layer_5 - estimator_1 - 5_)
[ 2019-04-23 22:01:25,466] [kfold_wrapper.log_eval_metrics] Accuracy(layer_5 - estimator_1 - 5_)
[ 2019-04-23 22:01:25,670] [kfold_wrapper.log_eval_metrics] Accuracy(layer_5 - estimator_1 - 5_)
[ 2019-04-23 22:01:25,671] [kfold_wrapper.log_eval_metrics] Accuracy(layer_5 - estimator_1 - 5_)

```

```

[ 2019-04-23 22:01:25,672] [cascade_classifier.calc_accuracy] Accuracy(layer_5 - train.classifi
[ 2019-04-23 22:01:25,845] [cascade_classifier.fit_transform] [layer=6] look_indexes=[0], X_cur
[ 2019-04-23 22:01:27,616] [kfold_wrapper.log_eval_metrics] Accuracy(layer_6 - estimator_0 - 5_
[ 2019-04-23 22:01:29,187] [kfold_wrapper.log_eval_metrics] Accuracy(layer_6 - estimator_0 - 5_
[ 2019-04-23 22:01:30,861] [kfold_wrapper.log_eval_metrics] Accuracy(layer_6 - estimator_0 - 5_
[ 2019-04-23 22:01:32,611] [kfold_wrapper.log_eval_metrics] Accuracy(layer_6 - estimator_0 - 5_
[ 2019-04-23 22:01:34,286] [kfold_wrapper.log_eval_metrics] Accuracy(layer_6 - estimator_0 - 5_
[ 2019-04-23 22:01:34,288] [kfold_wrapper.log_eval_metrics] Accuracy(layer_6 - estimator_0 - 5_
[ 2019-04-23 22:01:34,499] [kfold_wrapper.log_eval_metrics] Accuracy(layer_6 - estimator_1 - 5_
[ 2019-04-23 22:01:34,710] [kfold_wrapper.log_eval_metrics] Accuracy(layer_6 - estimator_1 - 5_
[ 2019-04-23 22:01:34,917] [kfold_wrapper.log_eval_metrics] Accuracy(layer_6 - estimator_1 - 5_
[ 2019-04-23 22:01:35,129] [kfold_wrapper.log_eval_metrics] Accuracy(layer_6 - estimator_1 - 5_
[ 2019-04-23 22:01:35,337] [kfold_wrapper.log_eval_metrics] Accuracy(layer_6 - estimator_1 - 5_
[ 2019-04-23 22:01:35,339] [kfold_wrapper.log_eval_metrics] Accuracy(layer_6 - estimator_1 - 5_
[ 2019-04-23 22:01:35,340] [cascade_classifier.calc_accuracy] Accuracy(layer_6 - train.classifi
[ 2019-04-23 22:01:35,514] [cascade_classifier.fit_transform] [layer=7] look_indexes=[0], X_cur
[ 2019-04-23 22:01:37,199] [kfold_wrapper.log_eval_metrics] Accuracy(layer_7 - estimator_0 - 5_
[ 2019-04-23 22:01:38,973] [kfold_wrapper.log_eval_metrics] Accuracy(layer_7 - estimator_0 - 5_
[ 2019-04-23 22:01:40,633] [kfold_wrapper.log_eval_metrics] Accuracy(layer_7 - estimator_0 - 5_
[ 2019-04-23 22:01:42,194] [kfold_wrapper.log_eval_metrics] Accuracy(layer_7 - estimator_0 - 5_
[ 2019-04-23 22:01:43,769] [kfold_wrapper.log_eval_metrics] Accuracy(layer_7 - estimator_0 - 5_
[ 2019-04-23 22:01:43,771] [kfold_wrapper.log_eval_metrics] Accuracy(layer_7 - estimator_0 - 5_
[ 2019-04-23 22:01:43,983] [kfold_wrapper.log_eval_metrics] Accuracy(layer_7 - estimator_1 - 5_
[ 2019-04-23 22:01:44,189] [kfold_wrapper.log_eval_metrics] Accuracy(layer_7 - estimator_1 - 5_
[ 2019-04-23 22:01:44,397] [kfold_wrapper.log_eval_metrics] Accuracy(layer_7 - estimator_1 - 5_
[ 2019-04-23 22:01:44,604] [kfold_wrapper.log_eval_metrics] Accuracy(layer_7 - estimator_1 - 5_
[ 2019-04-23 22:01:44,812] [kfold_wrapper.log_eval_metrics] Accuracy(layer_7 - estimator_1 - 5_
[ 2019-04-23 22:01:44,814] [kfold_wrapper.log_eval_metrics] Accuracy(layer_7 - estimator_1 - 5_
[ 2019-04-23 22:01:44,815] [cascade_classifier.calc_accuracy] Accuracy(layer_7 - train.classifi
[ 2019-04-23 22:01:44,982] [cascade_classifier.fit_transform] [layer=8] look_indexes=[0], X_cur
[ 2019-04-23 22:01:46,650] [kfold_wrapper.log_eval_metrics] Accuracy(layer_8 - estimator_0 - 5_
[ 2019-04-23 22:01:48,322] [kfold_wrapper.log_eval_metrics] Accuracy(layer_8 - estimator_0 - 5_
[ 2019-04-23 22:01:49,992] [kfold_wrapper.log_eval_metrics] Accuracy(layer_8 - estimator_0 - 5_
[ 2019-04-23 22:01:51,661] [kfold_wrapper.log_eval_metrics] Accuracy(layer_8 - estimator_0 - 5_
[ 2019-04-23 22:01:53,338] [kfold_wrapper.log_eval_metrics] Accuracy(layer_8 - estimator_0 - 5_
[ 2019-04-23 22:01:53,340] [kfold_wrapper.log_eval_metrics] Accuracy(layer_8 - estimator_0 - 5_
[ 2019-04-23 22:01:53,547] [kfold_wrapper.log_eval_metrics] Accuracy(layer_8 - estimator_1 - 5_
[ 2019-04-23 22:01:53,758] [kfold_wrapper.log_eval_metrics] Accuracy(layer_8 - estimator_1 - 5_
[ 2019-04-23 22:01:53,969] [kfold_wrapper.log_eval_metrics] Accuracy(layer_8 - estimator_1 - 5_
[ 2019-04-23 22:01:54,178] [kfold_wrapper.log_eval_metrics] Accuracy(layer_8 - estimator_1 - 5_
[ 2019-04-23 22:01:54,388] [kfold_wrapper.log_eval_metrics] Accuracy(layer_8 - estimator_1 - 5_
[ 2019-04-23 22:01:54,389] [kfold_wrapper.log_eval_metrics] Accuracy(layer_8 - estimator_1 - 5_
[ 2019-04-23 22:01:54,390] [cascade_classifier.calc_accuracy] Accuracy(layer_8 - train.classifi
[ 2019-04-23 22:01:54,562] [cascade_classifier.fit_transform] [layer=9] look_indexes=[0], X_cur
[ 2019-04-23 22:01:56,332] [kfold_wrapper.log_eval_metrics] Accuracy(layer_9 - estimator_0 - 5_
[ 2019-04-23 22:01:57,896] [kfold_wrapper.log_eval_metrics] Accuracy(layer_9 - estimator_0 - 5_
[ 2019-04-23 22:01:59,668] [kfold_wrapper.log_eval_metrics] Accuracy(layer_9 - estimator_0 - 5_
[ 2019-04-23 22:02:01,237] [kfold_wrapper.log_eval_metrics] Accuracy(layer_9 - estimator_0 - 5_

```

```
[ 2019-04-23 22:02:02,809] [kfold_wrapper.log_eval_metrics] Accuracy(layer_9 - estimator_0 - 5_)
[ 2019-04-23 22:02:02,813] [kfold_wrapper.log_eval_metrics] Accuracy(layer_9 - estimator_0 - 5_)
[ 2019-04-23 22:02:03,020] [kfold_wrapper.log_eval_metrics] Accuracy(layer_9 - estimator_1 - 5_)
[ 2019-04-23 22:02:03,227] [kfold_wrapper.log_eval_metrics] Accuracy(layer_9 - estimator_1 - 5_)
[ 2019-04-23 22:02:03,437] [kfold_wrapper.log_eval_metrics] Accuracy(layer_9 - estimator_1 - 5_)
[ 2019-04-23 22:02:03,643] [kfold_wrapper.log_eval_metrics] Accuracy(layer_9 - estimator_1 - 5_)
[ 2019-04-23 22:02:03,852] [kfold_wrapper.log_eval_metrics] Accuracy(layer_9 - estimator_1 - 5_)
[ 2019-04-23 22:02:03,853] [kfold_wrapper.log_eval_metrics] Accuracy(layer_9 - estimator_1 - 5_)
[ 2019-04-23 22:02:03,855] [cascade_classifier.calc_accuracy] Accuracy(layer_9 - train.classifi
[ 2019-04-23 22:02:03,856] [cascade_classifier.fit_transform] [Result][Reach Max Layer] opt_lay
```

```
In [46]: print('Training complete in {:.0f}m {:.4f}s'.format(
           time_elapsed // 60, time_elapsed % 60))
```

Training complete in 1m 41.7678s

```
In [47]: pred = predict_labels(gc,features_test.toarray())
         pred_scores.append(("DCF", [precision_score(labels_test,pred), recall_score(labels_test,
```

```
[ 2019-04-23 22:02:03,955] [cascade_classifier.transform] X_groups_test.shape=[[1672, 8710]]
[ 2019-04-23 22:02:04,064] [cascade_classifier.transform] group_dims=[8710]
[ 2019-04-23 22:02:04,065] [cascade_classifier.transform] X_test.shape=(1672, 8710)
[ 2019-04-23 22:02:04,131] [cascade_classifier.transform] [layer=0] look_indexes=[0], X_cur_test
[ 2019-04-23 22:02:04,983] [cascade_classifier.transform] [layer=1] look_indexes=[0], X_cur_test
[ 2019-04-23 22:02:05,795] [cascade_classifier.transform] [layer=2] look_indexes=[0], X_cur_test
[ 2019-04-23 22:02:06,605] [cascade_classifier.transform] [layer=3] look_indexes=[0], X_cur_test
[ 2019-04-23 22:02:07,421] [cascade_classifier.transform] [layer=4] look_indexes=[0], X_cur_test
[ 2019-04-23 22:02:08,230] [cascade_classifier.transform] [layer=5] look_indexes=[0], X_cur_test
[ 2019-04-23 22:02:09,037] [cascade_classifier.transform] [layer=6] look_indexes=[0], X_cur_test
[ 2019-04-23 22:02:09,847] [cascade_classifier.transform] [layer=7] look_indexes=[0], X_cur_test
[ 2019-04-23 22:02:10,654] [cascade_classifier.transform] [layer=8] look_indexes=[0], X_cur_test
[ 2019-04-23 22:02:11,456] [cascade_classifier.transform] [layer=9] look_indexes=[0], X_cur_test
```


```
In [48]: df = pd.DataFrame.from_items(pred_scores,orient='index', columns=['Precision', 'Recall',
           df
```

```
Out[48]:
```

	Precision	Recall	Accuracy	F1	Training Time (s)
SVC	0.990476	0.896552	0.984450	0.941176	0m 0.4593s
KN	1.000000	0.616379	0.946770	0.762667	0m 0.0041s
NB	0.982143	0.948276	0.990431	0.964912	0m 0.0018s
DT	0.900452	0.857759	0.967105	0.878587	0m 0.2270s
LR	0.933333	0.724138	0.954545	0.815534	0m 0.0107s
RF	1.000000	0.866379	0.981459	0.928406	0m 0.9180s
AdaBoost	0.952607	0.866379	0.975478	0.907449	0m 2.4091s
BgC	0.935185	0.870690	0.973684	0.901786	0m 1.1530s
ETC	0.995025	0.862069	0.980263	0.923788	0m 0.6672s

LSTM	0.854839	0.946429	0.971292	0.898305	0m 23.4645s
DCF	0.986842	0.969828	0.994019	0.978261	1m 41.7678s

```
In [49]: df.plot(kind='bar', y="Accuracy", ylim=(0.9,1.0), figsize=(11,6), align='center', col
plt.xticks(np.arange(11), df.index)
plt.ylabel('Accuracy Score')
plt.title('Distribution by Classifier')
plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)
plt.savefig("../img/sms-acc-v3.eps")
plt.show()
```



```
In [50]: import pickle
# dump
with open("../pkl/sms-gc-v3.pkl", "wb") as f:
    pickle.dump(gc, f, pickle.HIGHEST_PROTOCOL)

# # load
# with open("../pkl/2018_gc.pkl", "rb") as f:
#     gc = pickle.load(f)
```

0.08 Final verdict - gcForest is your friend in spam detection.

```
In [ ]:
```